

MOVING THE NEXT NOTES
to the iPad.

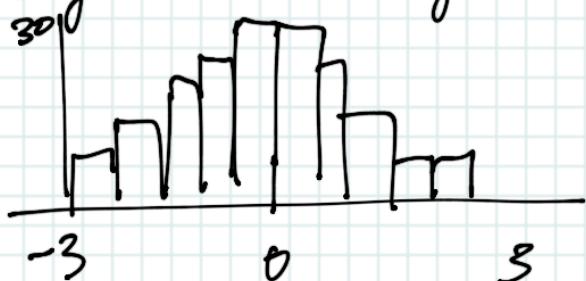
10 August 2020

So: trying to generate truth w/ more action in expectations (R, b).

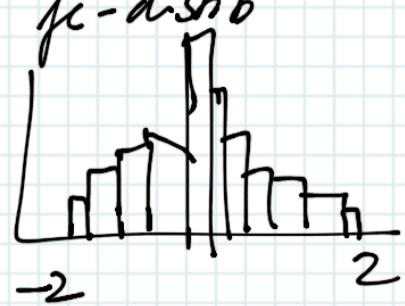
Observed:

- 1) If you increase the α -support, α 's in the simulation become more tight around zero, and extremes are smaller. (E.g. $\alpha \in (-2, 2)$ vs. $\alpha \in (-0.5, 0.5)$).
- 2) If you scale up α 's, α -distribution becomes more spread out: fatter tails

E.g. 4. α^{true} gives the following α -distrib



instead of
being like



and $\bar{\alpha}$ fluctuates much more.

Estimating $\hat{\alpha}$

- loss was initially much smaller 4.95 instead of on the order of 4000.
- estimation gets a lot of "fe was nan" messages and also a new one: "obj function returned NaN, trying a new point". Then...
 - ↳ I already got this message 9 times
- It took 585 sec.
- Implied k^{-1} is < 0 4 times!
 - maybe that's for negative k ?
- Again, $\text{loss}(\hat{\alpha}) = 3.2 < \text{loss}(\alpha^{\text{true}}) = 416$.
- Not possible to evaluate loss function?!
- ↳ It seems to be nans everywhere?
 - Is that why the loss just returned nan b/c all simulations were nan in the cross-section?
- Look at the loss: it sucks:
 - for α_1, α_5 , it's apparently always nan
 - for α_2, α_4 it blows up
 - for α_3 , it's oscillatory.

Any $\alpha > 0.5$ gives all nans. Maybe that's b/c since the truth is higher, it more quickly explodes?

- I'm not shocked that many histories explode.
- Also very many VAR instability problems.

$\hat{\alpha}$'s do more or less peak at the truth even now.
 (well, $\hat{\alpha}_3$ doesn't and $\hat{\alpha}_1$ & $\hat{\alpha}_5$ have 2 minima around the true value, but still)

→ Does this indicate that the data-generation is wrong or the estimation?

- True moments changed very much:
 - moments pertaining to x scaled up, ones pertaining to π or i scaled down / were unchanged
 - own moments or cross-moments w/ π or i have the same shape
 - cross-moments w/ x have changed shape:
 are now more U- or reverse U-shaped.

↳ I think these make sense b/c since κ is low, and $i = 1.5\pi$, x is the one that absorbs movements in expectations, while the pass-thru to π and to i is small.

⇒ The data-generation I think is correct. But this kinda suggests that the mean moments are also (more or

(loss) computed correctly, of which I'm not convinced.

↳ Let's try an estim w/ $\eta_x = 0.5$.

- But suppose the data is generated correctly. Why isn't the "x absorbs all" mechanism catchable by the estimation? I mean, moments didn't move far from initial ones.

↳ Explosion % reveals that in the estimation, either 0% or 100% of the N simulations explode. No, sometimes it's 1%, 79% or 99%. And it seems that the obj fn returns nan' when 100% explode. Yes.

I learn also, by the way, that each iter of the solver evaluates the loss function exactly 6 times.

→ Why? It's not changing α , or if it is, only one and only marginally.
I'm also noticing that it's only considering α 's really close to $\alpha_0 \Rightarrow$ multistab!

I'm also noticing that the simulation is not giving the 'if was nan' error most of the time, even though the solver gets a positive expl-percent. Hm!

Ask Ryan →

When the "fe was nan" error comes, then the explosion-counter does catch it. But apparently there are a bunch of explosions that come not from there.

Ok, it's clear: when fe is nan, it's always caught.
The other thing that causes "explosive" simulations is $k^{-1} < 0$. These get caught and register as explosions, but aren't really.

→ My check for "global nonnegativity" and wasn't catching everything!

Her ho : That's where the troubles start.

⇒ Initially, fe is nan (maybe the gain is too high) but then the bulk of the supposed explosions come from $k < 0$.

Ok... even when fe is nan it seems it's b/c or at the same time as $k < 0$...

→ It's at the same time. The initial iteration involves 5% "fe was nan" (5/100 simulations) but 100% have $k < 0$ at one point!

Let's see for the July 6 dataset ... and yes,

the initial iters involve 1% fcnan

69% $k < 0$

and later iters have 2% fcnan

86% $k < 0$

Converges after 23 iter.

Has no more problems after iter 4.

Why is fe often nan early in the search and not later? Is it b/c of the initial values x_0 ?

And could a finer, and more broad fe-grid for
the global nonnegativity test work?

! \hookrightarrow I introduced the variable broaden=2,
which extends the fegrid-file by \pm broaden at
both ends. It seems like this catches & subsumes
all the previous errors!

It still finds the same min though :S
(But at least it's 2/3 of the time)

If I make the loss fct=nan when $k < 0$ globally, I can
engineer the "Obj fct returned nan" error message.

Does this influence the estimate? (Previously I had set the objective to $1e+10$) No it doesn't.

→ I don't think so but it does affect the plots of the loss I would think!

Let's go back to the scaled-up truth and plot the loss. Should understand why $k < 0$: is it, as I think, that d isn't convex and so a f_e outside the grid is extrapolated to a $k < 0$?

The other thing is that this doesn't seem to really affect the behavior of the loss fit. It still converges to the same (wrong) thing. Actually, for the 4x4 truth, it converges to something different, but most likely that's b/c it took an avg of a smaller set.

The loss: $k < 0$ and $f_e = \text{nan}$ return. Does that mean that fgrid-fine needs to be over broader?

Not necessarily b/c now it happens that e.g.

$f_e = \text{nan}$ in 14%

but $k < 0$ in 0%.

In fact, $k < 0$ doesn't happen a lot! f_e is the problem here!

Why is it that for the estimation, k_{20} is the root of all evil, but in plotting the loss, suddenly $\text{je} = \text{van}$ becomes the problem #1? Is it that the estim avoids high α 's (or maybe also very low ones) that could cause turmoil?

Materials 40, Section 4: Look into behavior of simulation. I think I learn a lot from this. Fig 7, which shows means & distros of k^{-1} , $\bar{\pi}$ and je are indicative.

- 1) While k^{-1} spends a lot of time being near 0 (see histogram), $\text{mean}(k^{-1}) \approx 0.0185$ (see means)
↳ If the moments capture the latter aspect, then the estim should have a hard time pushing α_s down.
But the moments should more capture the first.
Moreover, what does it matter if $\text{mean}(k^{-1}) = 0.013$?
It still can enter the lower parts of its state space, and it does! So scrap that.
- 2) The cross-sectional mean of $\bar{\pi}$ & of $\text{je} \downarrow$ as $N \uparrow$.
Maybe this is why $\text{loss} \uparrow$ in N ?
I'm not sure this should happen! →

As $N \rightarrow \infty$, $\epsilon \rightarrow 0$. So $f_e \rightarrow 0$ and thus \bar{x} too.

But that means that mean moments will reflect a sim in which \bar{x} isn't moving much, and f_e are very small.

Now turn to changing α 's: do the following exercises:

- $\alpha = \alpha$
- More edges: $\alpha_1 \& \alpha_5$
- More middle: $\alpha_2 \& \alpha_4$
- More O-point: α_3

Obs. 1: If all α 's ≤ 0.1 , simulation doesn't blow up.

At 0.11, they do. (2 occurrences of "f_e was nan", one of which was also $k_{t-1} < 0$)

Obs. 2: $\alpha = 0$ doesn't cause explosions.

Obs 3: In the early scenarios, "f_e was nan" is the only message. In the late ones, it comes together w/ " $k_{t-1} < 0$ ". Does this mean that $k < 0$ when α_3 is large?

↳ Making the code output α suggests that:

- "f_e was nan" alone occurs when $\alpha_1 \& \alpha_5 \geq 0.11$
- "f_e was nan" + " $k_{t-1} < 0$ " occurs when $\alpha_2 \& \alpha_4 \geq 0.11$

- α_3 doesn't shift any water: it can also be 0.5, that doesn't cause explosions. (0-neighborhood indifference problem)

- these seem to occur across shock histories.

Now let's look at the plots of the means:

Scenario 1: Varying α_1 and α_3 : 0, 0.05, 0.1

- Shifts k^{-1} up from $(0.01, 0.02)$ to $(0.02, 0.03)$
- Shifts $\bar{\pi}$ from $(-0.01, 0.01)$ to $(-0.02, 0.04)$
- Barely shifts f_e : it remains in the $(-0.2, 0.2)$ range

Scenario 2: Varying α_2 and α_4

- Shifts k^{-1} up from $(0.005, 0.001)$ to $(0.05, 0.06)$
- Barely shifts $\bar{\pi}$: it remains in the $(-0.01, 0.03)$ range
- Barely shifts f_e : it remains in the $(-0.2, 0.2)$ range

Scenario 3: Varying α_3 :

- Shifts k^{-1} up from $(0.015, 0.02)$ to $(0.05, 0.06)$
- Barely shifts $\bar{\pi}$: it remains in the $(-0.01, 0.03)$ range
- Barely shifts f_e : it remains in the $(-0.2, 0.2)$ range

↳ So: why aren't $\alpha_{2,3,4}$ moving $\bar{\pi}$ & f_e ?

↳ is it b/c in that range k^{-1} . f_e is a really small number? i.e.: is the 0-neighborhood indifference a problem

here too?

→ I think so! I think what is happening is that the entire $\rho \in (-1, 1)$ -region yields too small \mathbf{E}^{-1} - ρ products such that they do matter a little bit (the loss does take a min there) but not a lot (the loss is very flat)

⇒ To be identified, you need to consider α 's associated w/ $\rho > 1$ (possibly even greater). But: you also need a truth that's not too big, b/c $\alpha > 0.1$ seems to be ill tolerated by the model.

→ To Do!

For Peter meeting:

- ① Loss does have a min, almost at right spot. (Fig 1)
 - a) Need convexity assumption - don't get removal (Fig 2)
 - b) Rescaling W affects loss, although no indication of inversion issues. (Fig 3.b)
 - c) Adding $E(\cdot)$ screws things up in a way I don't get. (Fig 4)
- (② Truth w/ more action in $E(\cdot)$) (Fig 6)
I thought that if $E(\cdot)$ aren't moving, then both rescaling

and informativeness of moments might be skewed up.
Explorations hinted at something else.
 $\alpha \stackrel{?}{\in} 0.1$ to avoid explorations.

③ Behavior of simulation

- Loss \uparrow in NP b/c $\bar{\pi}$ & fe \downarrow in N? (Fig 7 b & d)
- Fig 8: Only α_1 & α_5 can affect $\bar{\pi}$, and even that can't affect fe. (Fig 8)

Peter meeting

11 Aug 2020

Fig 1. Loss higher at $\hat{\alpha}$ vs α^{true}

Again: b/c data from model is nonlinear, this may not suggest that α^{true} is a local min.

Thinks that a lot of problems come from that (st a nonstat model is tricky).

The next issue: holding α 's fixed, but estimating one then one α may not hold. This would happen if you let $y = \beta_1 x_1 + \beta_2 x_2$ and x_1 & x_2 multicollinear.

\hookrightarrow might be some weird interaction between α 's so that the concavity restriction helps provide constraint.

Fig 8 conclusions sound exactly right.

Rescaling : still quite strange

| Fix α , check loss w/ and w/o scaling \rightarrow should recover the scaling factor.

Also explains why imposing convexity helps.

| Wann guard against claiming what is really a coding error. \rightarrow Should go back to that.

The level of loss fit has no meaning; only the curvature.

Min in Fig 1. don't line up w/ truth; not a problem b/c if you sim using a diff L , you'd get a diff min.

If std. errs large, can be b/c of

| ID \rightarrow doesn't go away as $N \rightarrow \infty$
Sampling error

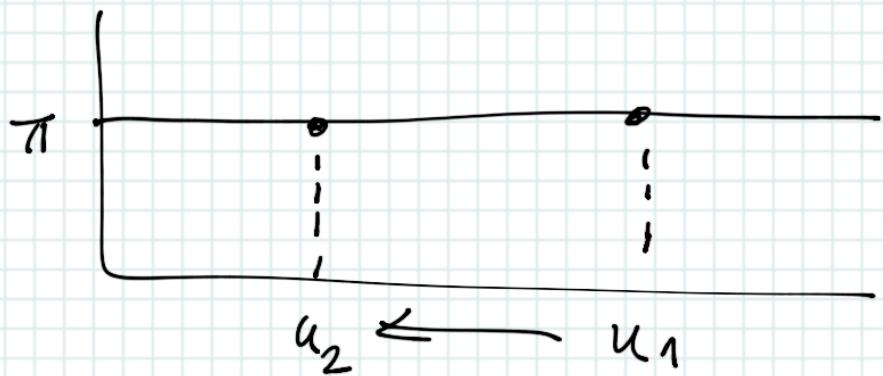
\hookrightarrow goes away as $N \rightarrow \infty$

Recall the logistic functional specification for anchoring fit:
Somehow more structure has to be imposed on
the anchoring function to guide the info in the
data

Analogy: Selma?
 Silvana Tenreyro: \hat{PC} : slope is $\approx 0 \rightarrow$ b/c CB is
 credible! So est a PC w/ these data doesn't
 tell you what would happen if CB abandoned
 target. (b/c that's just not in the data!)

\rightarrow bc obs that have modest fl, we don't learn a lot
 about how agents learn from forecast errors.

"flat PC": " π isn't moving in tandem w/ u anymore"



\hookrightarrow he says that the "flat PC" from refers to the idea that π doesn't follow u , not the other way around!

how u doesn't lead to any diff π !

- A fall-back approach: a simpler functional form for anch. pt.
- Or: pick unconditional moments w/o VAR or weighting matrix
 He called it an in-between when est & calibration.
 "Moment-matching": take sample moments & match those.

Successful esti will involve putting restrictions on α OR on the anchoring set to allow to glean info that is in the data.

↳ And I'm not sure why he calls the "moment-matching" by that name : it's not much different than my SMM.

Let's try the thing w/ α 's out in
the edges.

12 August 2020

Attempt (1) : July 6 data, $\text{fgrid} = [-2, -1.5, 0, 1.5, 2]$
(truth same) converged, not much change.

Attempt (2) : July 6 data, $\text{fgrid} = [-4, -3, 0, 3, 4]$

Took 500 sec & converged to sthg quite diff!
Not correct but maybe getting near?

Attempt (3) : July 6 data, $\text{fgrid} = [-4, -3 ; 3, 4]$ $wk=4$
Took 115 sec & converged to sthg very diff.

Isn't correct either but the thing is that the "truth" doesn't have a lot of fe in the $3-4.5$ range.

Attempt (4) : July 6 data, $\text{fgrid} = [-4, -3, 0, 3, 4]$
but w/ $W_{\text{mid}}=1000$.

Took 233 sec. Looks good. $fe \pm 3.84$ are still off, but I'm hopeful.

Attempt (5) : generate data with large forecast errors which you can est. w/ $\text{fgrid} = [-4, -3, 0, 3, 4]$

First try Aug 10 data ($g \cdot d^{\text{true}}$)

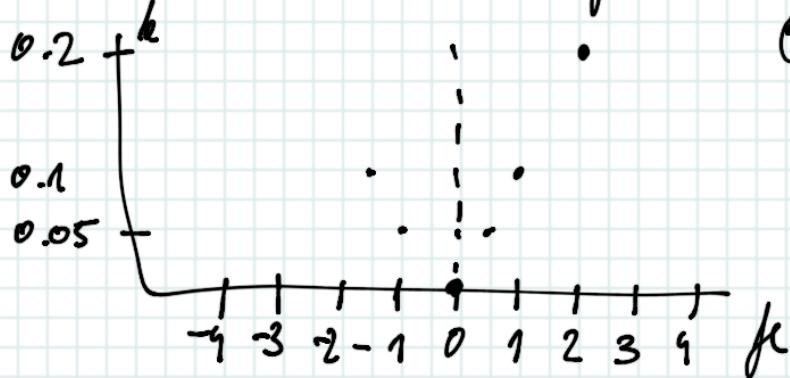
(w/ $W_{\text{mid}} = 1000$)

I'm also noting that the convexity things ain't working well.

Took 693 sec.

Got a bunch of "3% of sim histories exploded" messages.
There we go, lookin' good!

~~Attempt (6)~~: Make α 's associated w/ small f_c large, so the truth looks something like:



Oh shh. This is exactly the 4x times thing!

And $5 \times \alpha^{\text{true}}$ explodes already.

I don't think I can generate a dataset w/ bigger f_c b/c large f_c need large α 's, but sufficiently large α 's also explode, so...

Real data with Attempt (3) settings leads to a 1-2% sim histories exploded sign. 380 sec.

I think I'm going to leave it at that and consider the rescaling issue. Can I recover the scaling factor? And... no. Actually yes, if you recall that the loss involves squaring W .

↳ Ask Ryan: could it make a diff that in Ignoromia you specify the equations reside in W ?

Applied Young Economists Webinar (AYEW) 22 Aug 2020

Nanthu & Zexi Sun "Vague Talking at Central Banks' Press Conference: News or Noise?" (19 participants)

Hong Kong Institute for Monetary and Financial Research (note)

Vague talking that includes words like "risk", "uncertainty", "volatility" and "perturbation" raises stock returns
vs. → b/c investors expect expansionary trend in the future : the \oplus news dominate the \ominus news
that "current situation is bad".

Expansionary disinflation (Ball 1994)

dictionary from Loughran & ... , freq. of uncertain words

Regression: $R_t = b_0 + b_1 \Delta \text{vaguetalking}_t + \beta' b X_t + \epsilon_t$

First diff b/c EMH: stocks should only react to surprise

Prepare Ryan meeting:

- ① loss has a min at truth (Fig 1.)
- ② loss is indifferent to α 's such that $|f_e| \leq 1$ (Fig. 8)
- ③ Having a more action in fe truth and est- α 's corresponding to large fe works (Fig 15)

Details & problems:

- ④ Solver evaluates loss 6 times for the same α ?
- ⑤ Rescaling: changes shape of loss while no sign of inversion issue?

Ryan meeting

→ or play around w/ em.

12 August 2020

- Need to est 3 of shokes in order to get model-implied moments closer to those of the data, at least for π .
↳ If shokes too small, no $f_e \Rightarrow$ ID problems
- Isn't the whole story, b/c edgepoints also off
- Fig 8: would learn more if you plotted the auto-coranograms for these exercises!
→ So what's still amiss? We still don't know.

Weighting matrix: Eq (2): multiply by $\sqrt{\text{diag}(w)}$

→ Take the $\sqrt{\cdot}$ b/c otherwise you may put all weight on one moment.

• My rescaling strategy likely not to solve inversion issues b/c it's the ratio $\frac{\max(\text{element})}{\min(\text{element})}$

elements of Σ ^{which matters}. If it's more than 10^6 , then problem.

↳ Ryan's strategy
Fix largest el. of Σ and fix all elements to be not smaller than $1/10\ 000$ the largest.

• Not aligning the right els of W w/ the right moments could be a potential problem.

• Search alg:

6 evals b/c it computes gradient & jacobian.

(by 100 say) You can multiply the step for eval of ∇f , then the est. of derivatives is less precise but you get a better idea of the slope.

"first/second difference" or `diff` in Matlab

Search algorithm evaluates loss 6 times for each guess b/c if 1.) evaluates at the guess
2.) changes each x_i a LITTLE bit to estimate

first and second derivatives. It changes one α_i at a time, so this gives 5 additional coordinates.

Main point: the story that we can't use small β s to learn about how agents learn from β s is good, but it isn't the whole story. And one may be that shocks are too small, partly responsible for why there are no large β s in the data.

Work after

13 Aug 2020

I have 2 open businesses:

- 1) The ID question:
 - larger β needed in sample
 - 3 shocks
- 2) The weighting matrix:
 - take $\text{sqrt}(\cdot)$ of α s
 - figure out rescaling

Point 2): The weighting matrix.

Taking sqrt of the elements of W does change the est results. But in a sense these results are kinda better than the flat default Nsmmt.

For the "0's in the edges" specification, \sqrt{W} doesn't matter a lot. I guess not b/c that one is better ID-cd anyhow!

What I notice though is that taking \sqrt{w} , I don't get $L(\sqrt{w}) = L(w^{\text{rescale}}) \cdot \text{scaling factor}$
(that I don't get $L(w) = L(w^{\text{rescale}}) \cdot \text{scf}^2$ is no surprise)

Ok but what does hold is:

$$L(\sqrt{w}) = L(\sqrt{w^{\text{rescale}}}) \cdot \text{scf}. \quad \text{So it still works, as it should.}$$

The thing is: maybe, w/ larger shocks, $\text{Var}(E(\cdot))$ will be higher so the weighting matrix won't be an issue at all. Still: I need to understand why the weighting matrix rescaling changes the loss function.

Comparing to the default Nsimul picture, does \sqrt{w} change the shape of the loss? \rightarrow It does, a little, but not a lot. This is what I expected, given that it weights diff. moments and therefore also obtains diff. estimates.

Now scale w and see if the loss changes.

It does - but again only for x_1 & x_5 ...

I've also checked what Ryan said about the order being screwed up: that doesn't seem to be happening either.

there are 2 options: either I'm somehow rescaling in a way that does change relationships, or there is something else somewhere.

I think it's Option 2 b/c

$$\frac{\max(W)}{\min(W)} = \frac{\max(W^{\text{rescale}})}{\min(W^{\text{rescale}})}$$

also we should have

$$W^{\text{rescale}} = \frac{1}{\text{scf}} W \quad \text{and that holds too!}$$

↳ I'm really starting to think that the change in the shape of the loss fit has very weird to do w/ the indifference of $\alpha_2, \alpha_3, \alpha_4$.

The rescaling makes the loss smaller \rightarrow maybe the non-rescaled one would have the same shape if I zoomed in more?

Erm... not quite. Can it be that it's a numerical thingy that it changes shape?

Honestly, I don't think so.

Try to set $W = I$. Now rescale loss and see if it changes.
 \rightarrow it does! And the same way too!

I have an idea. Maybe the loss plot is changing shape b/c the convexity moment is in place.

Shut off $W_{\text{diffs2}} \rightarrow W_{\text{diffs2}} = 0$.

Set $W = I$ and compare loss w/ & w/o rescaling!
⇒ They have the same shape!

Redo w/ $W = \Sigma^{-1}$ & $W_{\text{diffs2}} = 0 \Rightarrow$ same thing!

YEAH! That was it! The convexity moment was effectively relatively weighted more when I scaled down the weights on the others!

I also don't get why evaluating the loss plot now takes longer than before. Producing these loss plots takes 35s instead of 75 sec. → today

→ Whoa, I can't believe that it was that simple!
So I can move on to what matters: the forecast errors and the variance of shocks. Let's think:
how best to do this:

A first pass thing is to see whether simulated f_t^A as \hat{f}_t^A . (command - check - simulation - approx. a)

Set $N=1$ ($\text{rng}(1)$). Default $b_2 = 1$ $\alpha = r, u, i$

Default: $f \in (-2, 2)$. $(\underline{-2, 2})$ $X = 2.5 \cdot \alpha^{\text{true}}$.

Raise $b_u = 2 \rightarrow \text{explodes}$. $\alpha = 4 \cdot \alpha^{\text{true}}$ $\alpha = \alpha^{\text{true}}$

Raise $b_u = 1.2 \Rightarrow f \in (-4, 4)$. $\text{explodes } (-2, 2)$

Raise $b_u = 1.5 \Rightarrow f \in (-5, 5)$. $-(-f) \in (-4, 4)$

$b_u = 4 \rightarrow (10, 10)$ $b_u = 1.8 \rightarrow \text{explodes}$. $-u \in (-5, 5)$

Raise $b_i = 1.2 \Rightarrow f \in (-2, 2) \setminus (-2, 2) \cap (-5, 5)$ ($b_i = 2$ too)

$b_i = 4 \Rightarrow \text{same. } (\underline{-2, 2}) \text{ explodes}$

b_i should have an effect though.

$b_i = 8 \Rightarrow f \in (-4, 4) \setminus (-2, 2)$ \rightarrow It does, just a much smaller one.

Raise $b_r = 8 \Rightarrow f \in (-5, 5)$ explodes , $b_r = 2$

\hookrightarrow Also a much smaller effect. I'm a little surprised that b_r apparently has a higher effect than b_i b/c ... oh no, idiot... of course b_r has a higher effect b/c the int. rate works from the IS-curve!

$$\pi = kX + u$$
$$X = b_i + r^n$$
$$i = \gamma_\pi \pi_f + \bar{i}$$

So the strength of the effect is of course $b_u > b_r > b_i$.

Ok, so that was like initial reconnoitering. I guess the next step is going to be to see how the estimation reacts to 1) the data being generated via higher β . Maybe the easiest is to do $\beta_0 = 4$ as a first pass. w/ $\alpha^{\text{true}} = \alpha^{\text{true}} (0.05, 0.025, 0, -11, -11)$

↳ This is the Aug 13 data!

Shit. VAR complaints, a lot of them.

Ridge? No complaints.

↳ But then maybe the estimation needs to use ridge too.

↳ Yep, it definitely complains.

Ok now it has some explosions (it was non) and it still has VAR issues despite the ridge. Let's try to lower (?) the ridge parameter λ .

(It was 0.001, now it's 0.0001)

↳ immediate VAR error.

So raise it to 0.01

↳ seemed to work well, but then an "imput to rank must not contain NaN or Inf" error. Hmmm.

↳ Let's try $\beta_0 = 2$. Think I need to rethink the ridge?

Yeah, immediate error. → Can create data w/o ridge? No.

Ok, so the situation is the following: for higher β , data creation and estimation needs ridge. So To Do: look into ridge and confirm that it's correct. Create some replicable simple example.

Confirming ridge

17 Aug. 2020

The issue of centering & scaling: Matlab's default for ridge is to center & scale: X is standardized (mean 0, var 1)

and y is centered (demeaned I guess?)

There is mention of this on some websites when I check ridge. Why is this important?

Statweb.Stanford.edu

Ridge is related to PCA: sort of weighted X 's are j^{th} principal components of X .

→ Ridge projects y on the PC's w/ large variance.

Statlect.com provides a hint as to why:

OLS is scale-invariant: β^{OLS} of $R \cdot X = R^{-1} \cdot \beta^{\text{OLS}}$ of X .

β^{ridge} is NOT scale-invariant! Therefore, whether you express variables in m or cm, e.g., matters!

↳ we therefore standardize all variables in ridge to

avoid undesirable scaling effects.

But I think you only need to stdize X , not Y ?

Matlab helpfile is confusing me now:

$$B0 = \left[\frac{\text{mean}(Y) - \text{mean}(X) \cdot \frac{B1}{\text{std}(X, 0, 1)}}{\text{std}(X, 0, 1)} , \frac{B1}{\text{std}(X, 0, 1)} \right]$$

Maybe Matlab is actually standardizing in each case,
it just computes $B0$ as above?

Stata notes on Ridge (ncss-wpengine.netdna-ssl.com)

- standardize all variables : X and Y
 - when final results are presented, they are adjusted back into their original scale. (Matlab's "0").
 - When "scaled=0" "ridge restores the coefficients to the scale of the original data".
- ↳ It's out:
- Matlab standardizes in any case
 - scale = 1 or 0 determines whether Matlab scales them back to their original scale.

→ So let's stdize Y & X !

Many observations:

1) still not the same as Matlab.

2) Matlab's ridge.m doesn't seem to be scaling Y .

↳ If I use the original, non-stdized Y , I get the same as Matlab's Ridge, scale = 1.
✓ Ok, but should one stdize Y or not?

Let's see:

$$\hat{\beta}^{\text{OLS}} = (X'X)^{-1} X'Y. \quad \text{Let } z = X R \xrightarrow{\text{rescale.}}$$

$$\begin{aligned}\tilde{\beta}^{\text{OLS}} &= ((X R)' X R)^{-1} (Z R)' Y \\ &= (R' X' X R)^{-1} R' X' Y \\ &= (R')^{-1} (X' X)^{-1} (R^{-1} R') X' Y \\ &= (R')^{-1} (X' X)^{-1} X' Y = R'^{-1} \hat{\beta}^{\text{OLS}}.\end{aligned}$$

$$\hat{\beta}^{\text{ridge}} = (X'X + \lambda I)^{-1} X'Y \quad \text{Again } z = X R$$

$$\begin{aligned}\tilde{\beta}^{\text{ridge}} &= ((X R)' X R + \lambda I)^{-1} (X R)' Y \\ &= (R' X' X R + \lambda I)^{-1} R' X' Y \\ &= (R' X' X R + \lambda R'(R')^{-1} R^{-1} R)^{-1} R' X' Y \quad \begin{array}{l} \text{Factor } R' \\ \text{and } R \end{array} \\ &= [R' (X' X + \lambda (R')^{-1} R^{-1}) R]^{-1} R' X' Y \\ &= R'^{-1} (X' X + \lambda (R')^{-1} R^{-1})^{-1} (R')^{-1} R' X' Y\end{aligned}$$

↑ inverse of product switching order.

$$\hat{\beta}^{\text{ridge}} = R^{-1} (X'X + \lambda (R')^{-1} R^{-1})^{-1} (R')^{-1} R' X' Y$$

$$= R^{-1} (X'X + \lambda \underline{(R')^{-1} R^{-1}})^{-1} X' Y$$

For this to be equal to $\hat{\beta}^{\text{ridge}}$, we need $(R')^{-1} R^{-1} = I$
i.e. R needs to be orthonormal. Generally doesn't hold!

↳ only need to standardize X ! (not Y)

Do need to standardize but also need to scale back! → 18 Aug 2020
NFA stats.stackexchange question & answer: (Philipp Burkhardt) For when only regressors std-
Let $z_j := \frac{x_j - \bar{x}_j}{s_j}$ standardized regressor j ^{ized.}

Then we have the regression line

$$E[Y] = \beta_0 + \sum_{j=1}^k \beta_j z_j$$

which gives fitted values

$$\hat{y} = \hat{\beta}_0 + \sum_{j=1}^k \hat{\beta}_j z_j . \quad \text{Write this as}$$

$$\hat{y} = \hat{\beta}_0 + \sum_{j=1}^k \hat{\beta}_j \left(\frac{x_j - \bar{x}_j}{s_j} \right)$$

Mallat has \hat{y} in her too.

$$\hat{y} = \left(\hat{\beta}_0 - \sum_{j=1}^k \hat{\beta}_j \frac{\bar{x}_j}{s_j} \right) + \sum_{j=1}^k \hat{\beta}_j \frac{x_j}{s_j}$$

Scaled-back intercept | Rearrange

$\hat{\beta}_j$ scaled back
($=$ what Mallat has)

If both regressors and regressed were scaled:

$$\hat{Y}_{\text{scaled}} = \frac{\hat{Y}_{\text{unscaled}} - \bar{Y}}{S_y} = \hat{\beta}_0 + \sum_{j=1}^k \hat{\beta}_j \left(\frac{x_j - \bar{x}_j}{S_j} \right) \quad | \cdot S_y \leftarrow \bar{y}$$

$$\begin{aligned}\hat{Y}_{\text{unscaled}} &= \bar{y} + S_y \hat{\beta}_0 + S_y \sum_{j=1}^k \hat{\beta}_j \left(\frac{x_j - \bar{x}_j}{S_j} \right) \\ &= \left(\bar{y} + S_y \hat{\beta}_0 - \underbrace{S_y \sum_{j=1}^k \hat{\beta}_j \bar{x}_j}_{\text{F}} \right) + \underbrace{S_y \sum_{j=1}^k \frac{\hat{\beta}_j}{S_j} x_j}_{\text{A}}\end{aligned}$$

This doesn't correspond to Matlab either: extra terms absent in Matlab.

I don't think Matlab does scaling Y . I suspect $\text{mean}(Y)$ is added back to take care of $E(Y)$ somehow?

Actually, Matlab's help for `ridge.m`, "Coefficient Scaling" explains it. Matlab standardizes X and centers Y (uses $\tilde{Y} := Y - \text{mean}(Y)$), which is why $\text{mean}(y)$ shows up in the intercept.

Ok: at least whether you use demeaned Y or not doesn't matter for β_1, \dots, β_k , only for β_0 , the intercept.

Hm! If you don't use demeaned Y , & scale back w/o \bar{y} , you get the same thing as if you used demeaned \bar{y} and scaled back w/ \bar{y} .

Ah - I think ridge is scale-invariant in Y :

$$\hat{\beta}^{\text{ridge}} = (X'X + \lambda I)^{-1} X'Y$$

Now let $Z = YR$

$$\hat{\beta}^{\text{ridge}}(Z) = (X'X + \lambda I)^{-1} X'ZR$$

$$= \hat{\beta}^{\text{ridge}} \cdot R$$

So just to be like Matlab, demean Y and standardize X .

Wait - why is $\hat{\beta}^{\text{ridge}}(Y) = \hat{\beta}^{\text{ridge}}(\underbrace{Y - \text{mean}(Y)}_{=: F})$?

$$\hat{\beta}_0^{\text{rescaled}} \underset{\substack{\uparrow \\ \text{0: intercept}}}{\text{adding back } \bar{Y}} = \hat{\beta}_0^{\text{rescaled}} \underset{\substack{\uparrow \\ \text{not}}}{} + F.$$

which makes sense. But this shouldn't be this entire business about scaling Y if it doesn't matter for $\hat{\beta}^{\text{ridge}}$...

Ah ... I think I know $X'Y$ is kinda $\text{Cov}(X, Y)$

$$\text{Cor}(X, Y) = \text{Cor}(X, Y - \bar{Y})$$

\uparrow b/c \bar{Y} is a constant ($\text{Var}(\bar{Y}) = 0$)

$$\text{Cov}(X, Y) = E[(X - E(X))(Y - E(Y))]$$

$$\text{Cov}(X, Y - \bar{Y}) = E[(X - E(X))(Y - \bar{Y} - E(Y - \bar{Y}))]$$

$$= E[(X - E(X))(Y - E(Y) - \bar{Y} + \bar{Y})] = \text{Cov}(X, Y).$$

→ So demeaning Y gives you the same $\hat{\beta}$ ridge as not demeaning.

So it seems like in any case you adjust the intercept, adding \bar{Y} to it. Ok...

Finally the last ridge - thing of interest:

$Y_{n \times k}$ (for VAR)

↳ Matlab can't do it, but doing it by hand, and then separately for each regressor gives the same thing.

↳ So more on: generate data w/ the new ridge commands and see if that solves the VAR-instability issues.

Note: the data-generation itself of 13 August ($\beta_n = 2$) complains of VAR-instability. That had the old ridge.

Data-generation complains w/ new ridge code?

(rf-var-ridge.m)

Nope, it didn't! :) ($\beta_n = 2, \lambda = 0.004$)

⇒ acf-sim-univariate-data-18-Aug-2020.

Ok, so finally do 1st w/ 18 Aug data

- ridge $\lambda = 0.001$

- $b_n = 2$

- \sqrt{W} (! this guy needs to become the new reference fig, as in Fig 3, Materials 41)

- $\gamma = 0.001$

Estimation: instability in VAR: 6 times

↳ It seems like it's the initial 6 evnts that cause problems. Warning never occurs afterwards.

→ Fig 4, Materials 41 ($\text{Flag} = 3$)

- Try same thing for $\text{Tolfun} = 1e-9 \rightarrow \text{same, Flag} = 2$.

- Same, $\text{TolX} = 1e-9 \rightarrow \text{same, Flag} = 3$.

- Now: do 'manual' grid out in far end ($\text{Tolfun} = \text{TolX} = e-9$)
But really one should have $\text{TolX} = \text{Tolfun} = 1e-6$ (default)
and add $\text{Wmid} = 1000$. Yeah this one's not great.

↳ Doing that now.

- tols back at defaults

- $b_n = 2$ (ridge, $\lambda = 0.001$)

- \sqrt{W}

- $\text{fgrid} = [-4, -3, 0, 3, 4]$

Not bad but not great either. What avenue to pursue? Raise b_n more?

Prep for Peter:

- ZEW interview
- Presentations : Aug 25 12.30 - 2pm (BU macro read)
- Sep 1 3.30 - 4.30 pm (Diss Workshop)
- Content: just Fig 15 from Materials 4D

Send link!
R

18 Aug 2020

Peter meeting

Main issue: works, but not so well in the neighborhood of 0 β .

↳ Ready to go back to true data? Yes?

But what you wanna do for sure is to constrain the parameters of β close to zero.

For real data: proceed in steps:

- first identify α (large β) and constrain α (small β) and see how well you can do
 - gradually expand # of α to expand to α (small β)
- Volatility of shocks important AND data is likely to be informative about them
- ↳ could est them
- ↳ or calib them such that model-sim data matches volatility of data.

• β shocks are simply harder to calibrate!

→ 2 questions:

1. In real world data, you wanna have a good feel for how ambitious you can be in est & and whether this depends on β
2. β is simply hard to assign values to.

↳ Talk after BN Seminar!

10:30 - 11:45 Tu & Th
FEB 9 - 20.15

Do an email exchange after DiSSWork to talk on Sep 2 or the next days.

Work after

OK - so see if we can estimate β .

I wanna do a new obj function which complements
obj - GMM - LGM gain - univariate - mean.m
in that it also estimates β_j $j = r, u, i$ + - shocks.m

I also wanna

✓ change names of figures so that they're not as long

→ Try it on Aug 10 data b/c that's the hx scaled-up truth.

! And TW is a default now! ! Want $= 1000$ too! Nsimul too!

The code is running... Took 713 sec. And it ends:

$$b_i = 0.1 \quad H_j \quad (\text{while truth is } 1). \quad \hookrightarrow 70 \text{ iter}$$

Why didn't it more? Error in code or? Is it time to try Ryan's thing of making it take larger steps when evaling the loss?

After 70 iterations, it looks like Matlab never tries anything else than b_0 . It doesn't seem like it's a problem of getting the values into "param".

→ I think it's something like "forward finite differences". As usual, some background: finite differences seems to be a numerical approximation method for derivatives.

Wait you idiot: of course it's not trying out any values of b : b is irrelevant in the model as it now stands. I think I need to modify the obj. fn.

→ Yeah, now I can see Matlab trying different things, that's good.

Finite difference := an expression of the form $f(x+b) - f(x+a)$
difference quotient := $\frac{f(x+b) - f(x+a)}{b-a}$ "finite diff approximation"

Finite differences are used in approximating derivatives, or for "finite difference methods" for solving ODE's.

finite difference quotients are used in approximating derivatives.

3 types :

$$\text{forward difference : } \Delta_h[f](x) := f(x+h) - f(x)$$

$$\text{backward difference : } \nabla_h[f](x) := f(x) - f(x-h)$$

$$\text{central difference : } S_h[f](x) := f(x + \frac{1}{2}h) - f(x - \frac{1}{2}h)$$

! Authors who use the term "finite differences" to mean "finite difference approximation of derivatives" also mean "forward difference quotient" when they say "forward difference".

Relation w/ derivatives

$$f'(x) := \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

For a fixed h , this is

$$\frac{f(x+h) - f(x)}{h} =: \frac{\Delta_h[f](x)}{h}$$

the forward difference quotient!

And I think the 45 poly correction to the convexity moment does this: $\Delta\alpha = f(x+h) - f(x)$, $h = \Delta$ fgrid.

Exciting stuff: in Matlab's solvers, there's the

Finite Difference Step Size: I guess this is " h ".

The default is $\sqrt{\text{eps}}$ for forward
(for optimset, "FinDiffRelStep") $\text{eps}^{(1/3)}$ for central
finite diff.

Finite Difference Type: "forward" (default) or "central"
Central is more accurate but takes twice
as many function evaluations.

(for optimset, "FinDiffType")

? 45 min!

I'm trying: same (i.e. MaxFunEvals = 200 (instead of 700))

but FinDiffStepSize = $100 \cdot \sqrt{\text{eps}}$.

Same. $\beta = (0.43, 0.25, 0.56)$.

Let's go back to MaxFunEvals = 700 (12 min).

and set FinDiffStepSize = $100\ 000 \cdot \sqrt{\text{eps}}$. ^{awful} _{timed out}

↪ Better but still timed out. Try $1M \cdot \sqrt{\text{eps}}$ ² didn't move

Let's make a little: why is the solver iterating out now
w/o convergence? Am I starting too far from the truth?
Is this complementarity between α and β which comes
from the \oplus feedback loop in the model?

What is really weird is that, compared to the case where

I fixed $\beta = 1$, now the estimates for $\hat{\alpha}$ are lower, while $\hat{\gamma} < \gamma^{\text{true}} = 1$. I would have expected $\hat{\alpha}$ to be higher, b/c the time data is more volatile ($b/c \gamma^{\text{true}} > \beta$) so to capture that, if the estimation doesn't raise β , it should raise $\hat{\alpha}$! Analogously, the moments aren't matched well at all.

Interestingly, increasing the `findiff stepsize` does allow the estimation to move away from the initial values more and fits moments better too. But why doesn't the solver continue on this path?

With `findiff stepsize` = $100k \cdot \sqrt{\epsilon}$, $\hat{\beta} = (3.5, 0.3, 0.75)$ and $\hat{\alpha} < \alpha^{\text{true}}$, but closer. So, in a sense, I see the tradeoff I was talking about (instead of $\hat{\alpha} \uparrow$, $\hat{\gamma} \uparrow$, in fact above γ^{true}), but in a way, $\hat{\beta} \uparrow$ was accompanied by $\hat{\alpha} \uparrow$, so the complementarity is kinda also not there.

I also wonder whether there's a kinda "symmetry indifference" in the sense that in half of simulations, $f_e < 0$, in the other half, $f_e > 0$, so $\alpha_{1,2} \gg 0$, $\alpha_{4,5} = 0$ matches moments just as well as $\alpha_{1,2} = 0$, $\alpha_{4,5} \gg 0$. Similarly, $\beta_r = 0$ & $\beta_i \gg 0$ may be undistinguishable from $\beta_r \gg 0$, $\beta_i = 0$.

Or is the estimation saying that x is more volatile and therefore $\hat{\beta}_r > \hat{\beta}_i$ or $\hat{\beta}_u$? But why isn't it saying $\hat{\alpha}$?

And is 1M · default pndiff stepsize not moving at all b/c

i) Jacobian is so imprecise

ii) or the loss is really flat over large surfaces?

$\frac{f(x+h) - f(x)}{h}$ is initially ≈ 0 as $\hat{\beta}$ is raised,

and then it becomes consistently $= \varepsilon \ll 0$, small.

Maybe what is going on, is that if $h :=$ pndiff stepsize is small, it's not getting to large $\hat{\beta}$, so it times out, but if h large, it gets to large $\hat{\beta}$, but then the steps in α are too large to make a diff.

Let's try $h = \begin{bmatrix} h_2 \\ h_\alpha \end{bmatrix} = \begin{bmatrix} \text{look for all the } \hat{\beta}'s \\ 1 \text{ for all the } \hat{\alpha}'s \end{bmatrix}$.

Interesting: this leads to $\hat{\alpha} < \hat{\alpha}^* \quad (h=100k \text{-default tparams})$
 $\hat{\beta} < \hat{\beta}^* \quad (-11-)$

Reverse to $h = \begin{bmatrix} h_2 \\ h_\alpha \end{bmatrix} = \begin{bmatrix} 1 \text{ for all } \hat{\beta}'s \\ \text{look for all } \hat{\alpha}'s \end{bmatrix}$

Interesting: this gets me almost as far as $h = 100k$ for all!
 \hookrightarrow So it's the h in α that matters, not so much the one in β .
 What does that mean? \rightarrow The loss is locally flatter in α than in β .

Try: $b = 100k \times \text{default}$ for all (uniformly), but initialize at truth: $b_0 = 1$.

↳ Much better (see Materials 41) \rightarrow seems to suggest that some of the "small step issue" may be coming from the initialization. But maybe the same holds true for α ?

↳ Trying $b_0 = 2$. Much more explosions and $k^{-1} < 0$.

But best so far! If I don't make the estimation "experience volatility", it doesn't.

↳ Maybe should do the same for α : initialize at truth or at higher values.

Also: I'm wondering if I should penalize $k^{-1} < 0$ in the estimation?

Prep for Ryan

- ZEW interview \rightarrow not avail.
- Presentations : Aug 25 12.30 - 2pm (BU macro read)
- Sep 1 3.30 - 4.30 pm (Diss workshop)
- Teaching: coordinate \rightarrow no rules or things to abide by.

Ryan meeting

Makes sense for $b_0 > b^{\text{true}}$.

19 August 2020

- Pick β_j such that you hit moments of data.

- Fixed-point problem: with a truth that looks like the data.

I fix some params. \rightarrow EST on real data. Take $\hat{\alpha}^{\text{real}}$ and put it into the fake data est. \rightarrow iterate until convergence.

capital crime
loss is

- Not continuous \rightarrow 2 cases you can make

• Kinney (misdemeanor)

\hookrightarrow cutting simulation that exploded introduces a jump in the obj. fn.

could explain
diff to Nestim

\hookrightarrow Any gain that's < 0 , set it to zero and

go on w/ the simulation

\hookrightarrow for exploding f_k , you could set f_k to a max

\hookrightarrow clean sol.

Worst-case scenario: is projection faculty which introduces kinks!

(can take 3×3 sigmas, est α for 9 cases, take $\arg(\hat{\alpha})$)
or the one w/ the best loss.

Try to calibrate β_0 so as to match moments of data. See if est then is better ID-ed. Can send him an email to get feedback.

Take-aways:

- ① 1st priority is to calibrate β 's to match data moments (for a given α) and see if this higher volatility would help ID-ing the α 's.
- ② Need to address how you deal w/ exploding $f_{\alpha \beta}$ in individual simulations. Throwing those out introduces jumps in the objective fn which is the biggest crime you can commit.
 - ↳ doing this could explain the difference to the "N estimations" strategy b/c that one doesn't have these jumps, it has a cross-section of $N=100$ anyway!
- ③ For the "ceteris paribus problem" in β and α , can start w/ a laboratory w/ fixed α, β . Estimate those for fake data. Does it work? Estimate on real data. Use $\hat{\alpha}$ & $\hat{\beta}$ as α^{true} & β^{true} in the next round of fake data. Iterate.